

Parallel in C (OpenMP)

Frederick Fung

National Computational Infrastructure, Australia

Acknowledgement of Country

Goal

This is a step-by-step introductory level OpenMP training session. The goal of this session is to get you familiar with some common practice of OpenMP and explore concurrency in for-loop.

- 1 The concepts
- 2 The coding exercises.

When do I need OpenMP?

When do I need OpenMP?

Among the choices of MPI, OpenCL, OpenAcc, CUDA, OpenSHMEM

OpenMP is an Application Program Interface (API) that accommodates easy multithread programming on shared memory.

OpenMP is an Application Program Interface (API) that accommodates easy multithread programming on shared memory.


This is ideal for programs running on single node of a cluster.

OpenMP is an Application Program Interface (API) that accommodates easy multithread programming on shared memory.


This is ideal for programs running on single node of a cluster.

For programs need to be parallel between multi-processors, consider MPI or hybrid MPI/OpenMP.

Threading: A schematic diagram



single-threaded process



multithreaded process

Figure 1: source: <https://www.cs.miami.edu/home/wuchtys/CSC322-21S/Content/UNIXProgramming/UNIXThreads.shtml>

Threads are NOT physical components of computers. A thread is a group of instructions to be executed by a core. Only one thread can run on a core at a time.

Threads are NOT physical components of computers. A thread is a group of instructions to be executed by a core. Only one thread can run on a core at a time.

When multithreads fork from a process, a core executes each thread for a short period of time and switch to other threads. Effectively, multiple threads are executed concurrently. (**Concurrency**)

Threads are NOT physical components of computers. A thread is a group of instructions to be executed by a core. Only one thread can run on a core at a time.

When multithreads fork from a process, a core executes each thread for a short period of time and switch to other threads. Effectively, multiple threads are executed concurrently. (**Concurrency**)

When multiple cores are used, the same number of threads can be executed in parallel. (**Parallelism**)




Figure 2: Three threads are forked from a process. They share most of the same execution resources in the process including a shared memory space[Pac11].

OpenMP is a **compiler directive** based API. To convert a serial program into a parallel program to run on a shared memory device, OpenMP uses compiler directives to tell the compiler to create a multithreaded version of the program. This limits the overhead of modifying codes to a minimal level for users.

Following closely to the principles in [MHK19], our hand-on session focus on using OpenMP to write multithreaded programs from three key ingredients:

- 1 Compiler Directives,
- 2 OpenMP Library,
- 3 Data Environment.

Example: Monte-Carlo Approximation of π

In the next part of our session, we use a Monte-Carlo method as a showcase to OpenMP programming.

Example: Monte-Carlo Approximation of π

In the next part of our session, we use a Monte-Carlo method as a showcase to OpenMP programming.

Generally speaking, Monte-Carlo method classifies a group of numerical algorithms that utilise random sampling to acquire numerical approximations.

Example: Monte-Carlo Approximation of π



In the next part of our session, we use a Monte-Carlo method as a showcase to OpenMP programming.

Generally speaking, Monte-Carlo method classifies a group of numerical algorithms that utilise random sampling to acquire numerical approximations.

The particular one that we are going to inspect is using the Monte-Carlo to approximate the value of π . Imagine scattering a handful of rice.

Example: Monte-carlo Approximation of π




Figure 3: Generate N random sampling points within a square, and count the number h of samples that fall in the circle. Then the approximation $\pi \approx 4h/N$

A serial code

```
seed  = 1; /* seed to generate random numbers */
for (i=0; i<N; i++){
    x = rand_r(&seed)/ (double) RAND_MAX; /* RAND_MAX to normalise */
    y = rand_r(&seed)/ (double) RAND_MAX;

    if (x*x + y*y <= 1.0f) count+=1;
}
```

x, y are normalised random numbers. Variable declarations are not shown.

Next we are going to walk through step by step to make this a parallel program.

Step 1: Parallel Construct

The parallel region of the code to be executed in multiple threads is specified by a **parallel construct** in OpenMP. A parallel construct is composed of the directive and the structured block.

```
seed = 1;
#pragma omp parallel[clause[,], clause...]
{ for (i=0; i<N; i++){
    x = rand_r(&seed)/ (double) RAND_MAX;
    y = rand_r(&seed)/ (double) RAND_MAX;

    if (x*x + y*y <= 1.0f) count+=1;
}
} // end of parallel region
```

Step 2: Library Routines

1

```
// request number of threads  
void omp_set_num_threads(int num_threads);
```

Step 2: Library Routines

```
1 // request number of threads
void omp_set_num_threads(int num_threads);

2 // return the thread rank in the team
int omp_get_thread_num();
```


Step 2: Library Routines

```
1 // request number of threads
  void omp_set_num_threads(int num_threads);

2 // return the thread rank in the team
  int omp_get_thread_num();

3 // total number of threads in the team
  int omp_get_num_threads();
```

Cyclic Distribution

Each thread works on loop iterations by increments determined by the rank of the thread.

```
#pragma omp parallel
{ int id = omp_get_thread_num();
  int numthreads = omp_get_num_threads();
  for (i=id; i<N; i+numthreads){
    seed = numthreads;
    x = rand_r(&seed) / (double) RAND_MAX;
    y = rand_r(&seed) / (double) RAND_MAX;

    if (x*x + y*y <= 1.0f) count+=1;
  }
} // end of parallel region
```

Block Distribution

Each threads is given an approximately equal sized block of consecutive loop iterations

```
#pragma omp parallel
{ int id = omp_get_thread_num();
  int numthreads = omp_get_num_threads();
  int start = id * N / numthread;
  int end = (id+1) * N / numthread;
  for (i=start; i<end; i++){
    seed = numthreads;
    x = rand_r(&seed) / (double) RAND_MAX;
    y = rand_r(&seed) / (double) RAND_MAX;
    if (x*x + y*y <= 1.0f) count+=1;
  }
} // end of parallel region
```

SPMD (Single Program Multiple Data)

Both of the two methods we discussed about use the SPMD pattern, i.e., the programmer has to define the copies of data for each thread.

SPMD (Single Program Multiple Data)

Both of the two methods we discussed about use the SPMD pattern, i.e., the programmer has to define the copies of data for each thread.

This can be error-prone since it modifies the code significantly by introducing new variables for each thread.

SPMD (Single Program Multiple Data)

Both of the two methods we discussed about use the SPMD pattern, i.e., the programmer has to define the copies of data for each thread.

This can be error-prone since it modifies the code significantly by introducing new variables for each thread.

Therefore, the two methodologies are difficult to implement. We let OpenMP to handle this for us

Step 3: Worksharing-loop Construct

To avoid SPMD pattern, OpenMP uses **worksharing-loop construct** to achieve loop-level parallelism. In essence, it leaves the compiler to split the loop iterations and the user does not need to modify the for-loop.

Step 3: Worksharing-loop Construct

To avoid SPMD pattern, OpenMP uses **worksharing-loop construct** to achieve loop-level parallelism. In essence, it leaves the compiler to split the loop iterations and the user does not need to modify the for-loop.

```
#pragma omp parallel
{seed = omp_get_thread_num();
#pragma omp for
for (i=0; i<N; i++){
    x = rand_r(&seed) / (double) RAND_MAX;
    y = rand_r(&seed) / (double) RAND_MAX;
    if (x*x + y*y <= 1.0f) count+=1;
}
}
```


If the Parallel construct is immediately followed by a worksharing-loop construct, then both directive can be combined. Hence, the following two patterns are equivalent.

```
#pragma omp parallel  
{  
  #pragma omp for  
    //for loop  
}
```

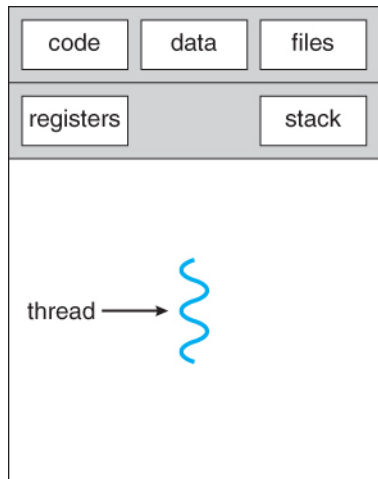
```
#pragma omp parallel for  
  //for loop
```

Race Condition


When a variable is shared between multiple threads, multiple threads can potentially access the same memory location concurrently. Then the update has unpredictable behaviour. It is seen in the **count** variable in our example.

```
#pragma omp parallel
{
    /* skipped */
#pragma omp for
    /* skipped */
    if (x*x + y*y <= 1.0f) count+=1;
}
```

Threading: A schematic diagram



single-threaded process



multithreaded process

Figure 4: source: <https://www.cs.miami.edu/home/wuchtys/CSC322-21S/Content/UNIXProgramming/UNIXThreads.shtml>

Time	Thread A	count	Thread B
T1	Read	i	
T2		i	Read
T3	Add	i	
T4		i	Add
T5	Write	i+1	
T6		i +1	Write

Table 1: Two thread update a value at the same memory location.

Step 4.5: Synchronisation

A **critical construct** enforces the attached block of code to be executed with mutual exclusion.

```
#pragma omp parallel
{
    /* skipped */
    #pragma omp for
    /* skipped */
    if (x*x + y*y <= 1.0f)
    {
        #pragma omp critical
        count+=1;
    }
}
```

Time	Thread A	count	Thread B
T1	Read	i	
T2	Add	i	
T3	Write	i+1	

Table 2: Thread A arrives the critical region first, thread A gets to execute the critical construct exclusively.

Critical

Although using critical construct renders the thread safety, it increases the serial portion. In some scenarios, if the critical block takes up a considerable amount of time then the code essentially gets serialised.

Critical

Although using critical construct renders the thread safety, it increases the serial portion. In some scenarios, if the critical block takes up a considerable amount of time then the code essentially gets serialised.

We have a better way to avoid using the expensive synchronisation in this case by smartly allocate data storage attributes. (We will come back on this)

Atomic

Similar to the critical construct, the **atomic construct** occurs with the mutual exclusion. Due to the different implementation, atomic has less overhead but the trade-off is that it also has more specific use cases.

Atomic

Similar to the critical construct, the **atomic construct** occurs with the mutual exclusion. Due to the different implementation, atomic has less overhead but the trade-off is that it also has more specific use cases.

More precisely, atomic only guarantees the mutual exclusion for operations that directly involves the storage location in memory.

```
#pragma omp atomic  
count+=1; /* thread safety protected */
```

```
#pragma omp atomic  
count += calc_count() /* potential data race */
```

```
#pragma omp critical  
count += calc_count() /* thread safety protected */
```

Barrier

A barrier defines a point at which all threads must arrive before any thread may proceed past the barrier.

Barrier

A barrier defines a point at which all threads must arrive before any thread may proceed past the barrier.

```
#pragma omp barrier
```

Barrier

A barrier defines a point at which all threads must arrive before any thread may proceed past the barrier.

```
#pragma omp barrier
```

We've been implicitly using barriers! Every worksharing construct has a barrier at the end to guard the thread safety.




Figure 5: Three threads are forked from a process. A barrier is places before they join as one master thread.[Pac11]

Nowait

A barrier is very expensive, the waiting time depends on the slowest thread. The **nowait construct** omits a barrier if it is determined not needed.

```
#pragma omp for nowait /* omits the barrier for  
the worksharing-loopF */
```

Nowait

A barrier is very expensive, the waiting time depends on the slowest thread. The **nowait construct** omits a barrier if it is determined not needed.

```
#pragma omp for nowait /* omits the barrier for  
                        the worksharing-loopF */
```

```
#pragma omp parallel  
{  
  #pragma omp for  
    { /* for loop */ } /* compiler places nowait */  
}
```

Recall that we used critical construct (or atomic) to avoid the data race from updating the **count** variable. However, this serialises the multithreaded code.

Recall that we used critical construct (or atomic) to avoid the data race from updating the **count** variable. However, this serialises the multithreaded code.

The race condition raises when multiple threads try to write to the same memory location. This condition will not present itself if every thread has its private **count** variable.

Recall that we used critical construct (or atomic) to avoid the data race from updating the **count** variable. However, this serialises the multithreaded code.

The race condition raises when multiple threads try to write to the same memory location. This condition will not present itself if every thread has its private **count** variable.

Reduction

Creates copy of a scalar variable local to each thread. Uses the clause operator at end of the worksharing-construct to combine all local copies into a shared global variable.

Step 5: Reduce the Scalar

Reduction

```
reduction (operation:variable)
```

```
#pragma omp parallel
{
    /* not shown */
#pragma omp for reduction(+:count)
    /* for loop not shown */
    if (x*x + y*y <= 1.0f) count+=1;
}
```

Exercise: Add parallel construct and worksharing-loop construct to turn the Monte-Carlo π program into parallel.

Step 6: Storage Attribute

Shared

A variable is **shared** if it only has one global copy and can be accessed by all threads in the team.

```
#pragma omp parallel shared(list)
```

Private

The **private** clause tells the compiler to create a new variable for the same type and the same name for each thread. The private variable is uninitialised in each thread.

```
#pragma omp parallel private(list)
```

Firstprivate

Like private clause, **firstprivate** tells compiler to create private copy for each thread. The only difference is that with firstprivate, the new private variable is initialised by copying the value of the original variable.

```
#pragma omp parallel firstprivate(list)
```


What variables in Monte-Carlo π should be created as private? We didn't add the private clause and why it still worked?

```
#pragma omp parallel
{seed = omp_get_thread_num();
#pragma omp for
for (i=0; i<N; i++){
    x = rand_r(&seed) / (double) RAND_MAX;
    y = rand_r(&seed) / (double) RAND_MAX;
    if (x*x + y*y <= 1.0f) count+=1;
}
}
```

Dafualt(none)

Data storage attribute is a very error-prone component of OpenMP programming (as we see in the Monte-Carlo π), use `default(none)` to force yourself to explicitly list the correct storage attributes on variables.

```
#pragma omp parallel default(none)
```

Ordered

The **ordered** enforces the execution to run sequentially in the same order as if it was executed in serial.

```
#pragma omp ordered
```

Step 7: Scheduling

The default worksharing-loop construct leaves the compiler to decide how to split a loop.

Step 7: Scheduling

The default worksharing-loop construct leaves the compiler to decide how to split a loop.

Recall we mentioned two parallel methodologies and they are equivalent in a sense that every thread gets an equivalent amount of work if loop iterations repeat the same execution. It is the case in our Monte-Carlo π exercise.

Static Schedule

Static means that the loop schedule is managed at compile time. The static schedule has a chunk parameter. When the chunk is not specified, the loop is split as described in the second methodology (block distribution). If the chunk is given, it is similar to the second methodology (cyclic distribution).

Static Schedule

Static means that the loop schedule is managed at compile time. The static schedule has a chunk parameter. When the chunk is not specified, the loop is split as described in the second methodology (block distribution). If the chunk is given, it is similar to the second methodology (cyclic distribution).

```
#pragma omp for schedule(static, chunk)
```

As a programmer, we know better than the compiler about the execution of our codes. Can we schedule the loop that is more efficient?

As a programmer, we know better than the compiler about the execution of our codes. Can we schedule the loop that is more efficient?

Dynamic Schedule

If we know the work for each loop iteration is a variable and can only be determined at the runtime, we may let the OS to decide the assignment of iterations to the threads. (Be careful, it comes with the cost of overheads.)

```
#pragma omp for schedule(dynamic, chunk)
```

In the next exercise, we will use most of the constructs and clauses that we've learnt today to concurrently generate datasets and visualise some pretty fractal pictures!

Definition

The Mandelbrot set is the set of complex numbers c for which the function $f_c(z_{i+1}) = z_i^2 + c$ does not diverge to infinity when iterated from $z_0 = 0$

$$M := \{c \in \mathbb{C} \mid z_{i+1} = z_i^2 + c < \infty, z_0 = 0, \quad i = 0, 1, \dots, N\}$$

Fractal Structure

The image of Mandelbrot Set exhibits the self-similar fractal structure. It is also an infinite set, more points means higher resolution and we can see more fine-grained details of the fractal image!




Figure 6: Mandelbrot set

Mandelbrot Set




Figure 7: Mandelbrot set zoomed in

Task: Add a parallel construct and appropriate clauses to generate the Mandelbrot set in parallel.

Things to think about:

- 1 What storage attributes should be given to the variables?

Task: Add a parallel construct and appropriate clauses to generate the Mandelbrot set in parallel.

Things to think about:

- 1 What storage attributes should be given to the variables?
- 2 What schedule should be assigned to the for-loop?

Task: Add a parallel construct and appropriate clauses to generate the Mandelbrot set in parallel.

Things to think about:

- 1 What storage attributes should be given to the variables?
- 2 What schedule should be assigned to the for-loop?
- 3 Which part of the code needs to be synchronised?

Task: Add a parallel construct and appropriate clauses to generate the Mandelbrot set in parallel.

Things to think about:

- 1 What storage attributes should be given to the variables?
- 2 What schedule should be assigned to the for-loop?
- 3 Which part of the code needs to be synchronised?

The CG method is a numerical iterative direct solver for solving linear system:

$$Ax = b, \tag{1}$$



where A is a symmetric positive definite matrix.

Algorithm 1: Standard CG (A, b , tolerance)

Result: x

```
1 Compute  $r_0 := b - Ax_0, p_0 := r_0$ ;  
2 for  $i = 0, 1, \dots, \dim(A)$  do  
3    $\alpha_i := (r_i, r_i) / (Ap_i, p_i)$ ;  
4    $x_{i+1} := x_i + \alpha_i p_i$ ;  
5    $r_{i+1} := r_i - \alpha_i Ap_i$ ;  
6   if  $r_{i+1} < \text{tolerance}$  then  
7     break  
8   end  
9    $\beta_i := (r_{i+1}, r_{i+1}) / (r_i, r_i)$ ;  
0    $p_{i+1} := r_{i+1} + \beta_i p_i$   
1 end
```

Task: Add a parallel construct and appropriate clauses to accelerate iterations of CG method.

-  Timothy G Mattson, Yun Helen He, and Alice E Koniges, *The openmp common core: Making openmp simple again*, MIT Press, 2019.
-  Peter Pacheco, *An introduction to parallel programming*, Elsevier, 2011.

Thank you!

Please help us improve!

https:
//anu.au1.qualtrics.com/jfe/form/SV_0BRRQvSz5oiwOqO